

Notebook

March 19, 2024

1 Problem Statement

A bike-sharing system is a service in which bikes are made available for shared use to individuals on a short term basis for a price or free. Many bike share systems allow people to borrow a bike from a “dock” which is usually computer-controlled wherein the user enters the payment information, and the system unlocks it. This bike can then be returned to another dock belonging to the same system.

A US bike-sharing provider BoomBikes has recently suffered considerable dips in their revenues due to the ongoing Corona pandemic. The company is finding it very difficult to sustain in the current market scenario. So, it has decided to come up with a mindful business plan to be able to accelerate its revenue as soon as the ongoing lockdown comes to an end, and the economy restores to a healthy state.

In such an attempt, BoomBikes aspires to understand the demand for shared bikes among the people after this ongoing quarantine situation ends across the nation due to Covid-19. They have planned this to prepare themselves to cater to the people’s needs once the situation gets better all around and stand out from other service providers and make huge profits.

They have contracted a consulting company to understand the factors on which the demand for these shared bikes depends. Specifically, they want to understand the factors affecting the demand for these shared bikes in the American market. The company wants to know:

Which variables are significant in predicting the demand for shared bikes. How well those variables describe the bike demands Based on various meteorological surveys and people’s styles, the service provider firm has gathered a large dataset on daily bike demands across the American market based on some factors.

2 Business Goal

You are required to model the demand for shared bikes with the available independent variables. It will be used by the management to understand how exactly the demands vary with different features. They can accordingly manipulate the business strategy to meet the demand levels and meet the customer’s expectations. Further, the model will be a good way for management to understand the demand dynamics of a new market.

```
[45]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

```

import seaborn as sns
import statsmodels
import statsmodels.api as sm
import sklearn
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.preprocessing import MinMaxScaler
from statsmodels.stats.outliers_influence import variance_inflation_factor

import warnings
warnings.filterwarnings('ignore')

```

```

[46]: bike_sharing_data = pd.read_csv('day.csv')      ## ../15 Upgrad/05 Machine_
      ↪ Learning 01/Assignment-Bike-Sharing/
      bike_sharing_data.head()

```

```

[46]:
   instant  dteday  season  yr  mnth  holiday  weekday  workingday  \
0         1  01-01-2018     1   0     1         0         6           0
1         2  02-01-2018     1   0     1         0         0           0
2         3  03-01-2018     1   0     1         0         1           1
3         4  04-01-2018     1   0     1         0         2           1
4         5  05-01-2018     1   0     1         0         3           1

   weathersit  temp  atemp  hum  windspeed  casual  registered  \
0         2  14.110847  18.18125  80.5833  10.749882    331      654
1         2  14.902598  17.68695  69.6087  16.652113    131      670
2         1   8.050924   9.47025  43.7273  16.636703    120     1229
3         1   8.200000  10.60610  59.0435  10.739832    108     1454
4         1   9.305237  11.46350  43.6957  12.522300     82     1518

   cnt
0   985
1   801
2  1349
3  1562
4  1600

```

```

[47]: bike_sharing_data.shape      # data has 730 records and 16 columns (or
      ↪ features)

```

```

[47]: (730, 16)

```

3 Finding the datatypes of the columns in the dataset

```

[48]: bike_sharing_data.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 730 entries, 0 to 729
Data columns (total 16 columns):
#   Column          Non-Null Count  Dtype
---  -
0   instant         730 non-null   int64
1   dteday          730 non-null   object
2   season          730 non-null   int64
3   yr              730 non-null   int64
4   mnth            730 non-null   int64
5   holiday         730 non-null   int64
6   weekday         730 non-null   int64
7   workingday      730 non-null   int64
8   weathersit       730 non-null   int64
9   temp            730 non-null   float64
10  atemp           730 non-null   float64
11  hum             730 non-null   float64
12  windspeed       730 non-null   float64
13  casual          730 non-null   int64
14  registered       730 non-null   int64
15  cnt             730 non-null   int64
dtypes: float64(4), int64(11), object(1)
memory usage: 91.4+ KB

```

4 No Null/blank values obseved in the dataset

```
[49]: bike_sharing_data.isnull().sum()
```

```

[49]: instant         0
      dteday          0
      season          0
      yr             0
      mnth           0
      holiday         0
      weekday         0
      workingday      0
      weathersit       0
      temp            0
      atemp           0
      hum             0
      windspeed       0
      casual          0
      registered       0
      cnt             0
      dtype: int64

```

5 checking for duplicate records and no duplicates are seen

```
[50]: duplicateRows = bike_sharing_data[bike_sharing_data.duplicated()]
      duplicateRows.count()
```

```
[50]: instant      0
      dteday      0
      season      0
      yr          0
      mnth        0
      holiday     0
      weekday     0
      workingday  0
      weathersit   0
      temp        0
      atemp       0
      hum         0
      windspeed   0
      casual      0
      registered  0
      cnt         0
      dtype: int64
```

6 Removing unwanted columns, viewing left over columns

```
[51]: bike_sharing_data = bike_sharing_data.drop(['instant', 'dteday'], axis=1)
      bike_sharing_data.columns
```

```
[51]: Index(['season', 'yr', 'mnth', 'holiday', 'weekday', 'workingday',
          'weathersit', 'temp', 'atemp', 'hum', 'windspeed', 'casual',
          'registered', 'cnt'],
          dtype='object')
```

7 ‘casual’ + ‘registered’ = ‘cnt’. Target var is ‘cnt’, so removing columns ‘casual’ and ‘registered’

```
[52]: bike_sharing_data = bike_sharing_data.drop(['casual', 'registered'], axis=1)
      bike_sharing_data.columns
```


```
[52]: Index(['season', 'yr', 'mnth', 'holiday', 'weekday', 'workingday',
          'weathersit', 'temp', 'atemp', 'hum', 'windspeed', 'cnt'],
          dtype='object')
```

8 converting 'yr' to int64 from categorical

```
[53]: bike_sharing_data['yr'] = bike_sharing_data.yr.astype('int64')
```

9 Boxplot for 4 important numeric columns

```
[54]: df = pd.DataFrame(data = bike_sharing_data, columns = ['temp', 'atemp', 'hum', 'windspeed'])
sns.boxplot(x="variable", y="value", data=pd.melt(df))
plt.show()
```



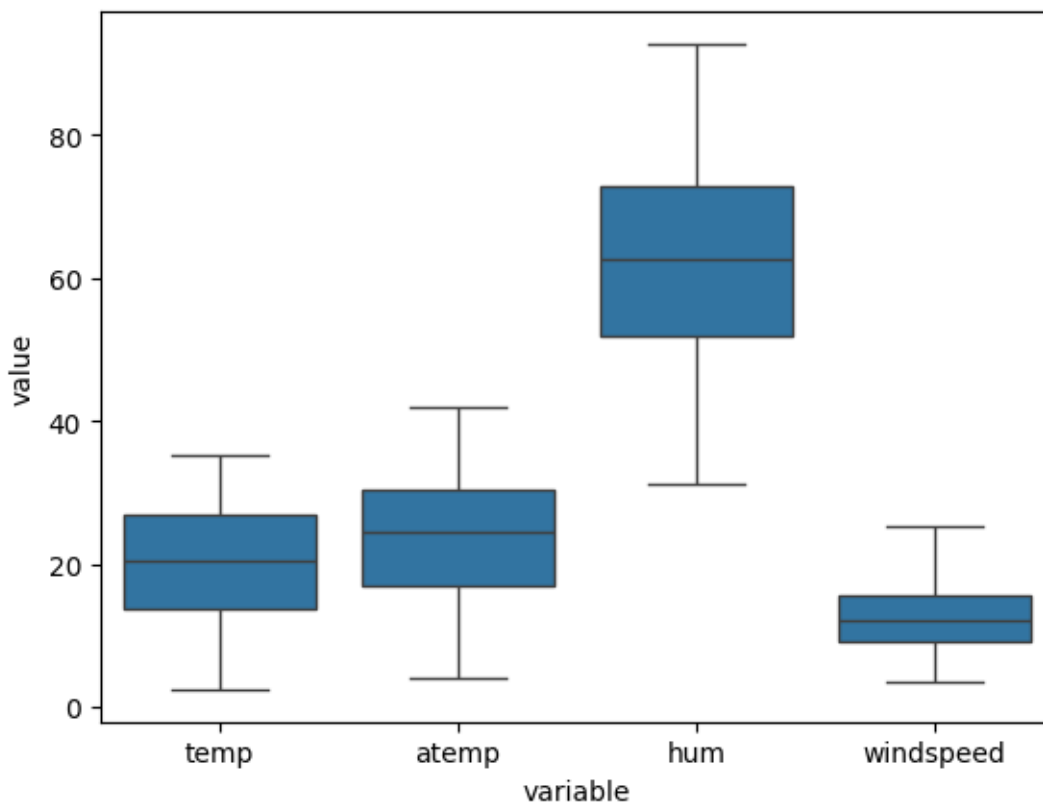
10 Treating outlier for columns 'hum' and 'windspeed'

```
[55]: upper_limit = bike_sharing_data['hum'].quantile(0.99)
lower_limit = bike_sharing_data['hum'].quantile(0.01)
bike_sharing_data['hum'] = np.where(bike_sharing_data['hum'] >= upper_limit,
    ↳ upper_limit, np.where(bike_sharing_data['hum'] <= lower_limit, lower_limit,
    ↳ bike_sharing_data['hum']))
```

```
[56]: upper_limit = bike_sharing_data['windspeed'].quantile(0.98)
lower_limit = bike_sharing_data['windspeed'].quantile(0.01)
bike_sharing_data['windspeed'] = np.where(bike_sharing_data['windspeed'] >=
↳upper_limit, upper_limit, np.where(bike_sharing_data['windspeed'] <=
↳lower_limit, lower_limit, bike_sharing_data['windspeed']))
```

11 Checking the data distribution through boxplots again, after treating the outliers

```
[57]: df = pd.DataFrame(data = bike_sharing_data, columns =
↳['temp', 'atemp', 'hum', 'windspeed'])
sns.boxplot(x="variable", y="value", data=pd.melt(df))
plt.show()
```



12 changing numeric columns to categorical

```
[58]: bike_sharing_data['season'] = bike_sharing_data.season.astype('category')
bike_sharing_data['weathersit'] = bike_sharing_data.weathersit.
      ↪astype('category')
bike_sharing_data['mnth'] = bike_sharing_data.mnth.astype('category')
bike_sharing_data['weekday'] = bike_sharing_data.weekday.astype('category')
```

```
[59]: bike_sharing_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 730 entries, 0 to 729
Data columns (total 12 columns):
#   Column      Non-Null Count  Dtype
---  -
0   season      730 non-null   category
1   yr          730 non-null   int64
2   mnth        730 non-null   category
3   holiday     730 non-null   int64
4   weekday     730 non-null   category
5   workingday  730 non-null   int64
6   weathersit   730 non-null   category
7   temp        730 non-null   float64
8   atemp       730 non-null   float64
9   hum         730 non-null   float64
10  windspeed   730 non-null   float64
11  cnt         730 non-null   int64
dtypes: category(4), float64(4), int64(4)
memory usage: 49.7 KB
```

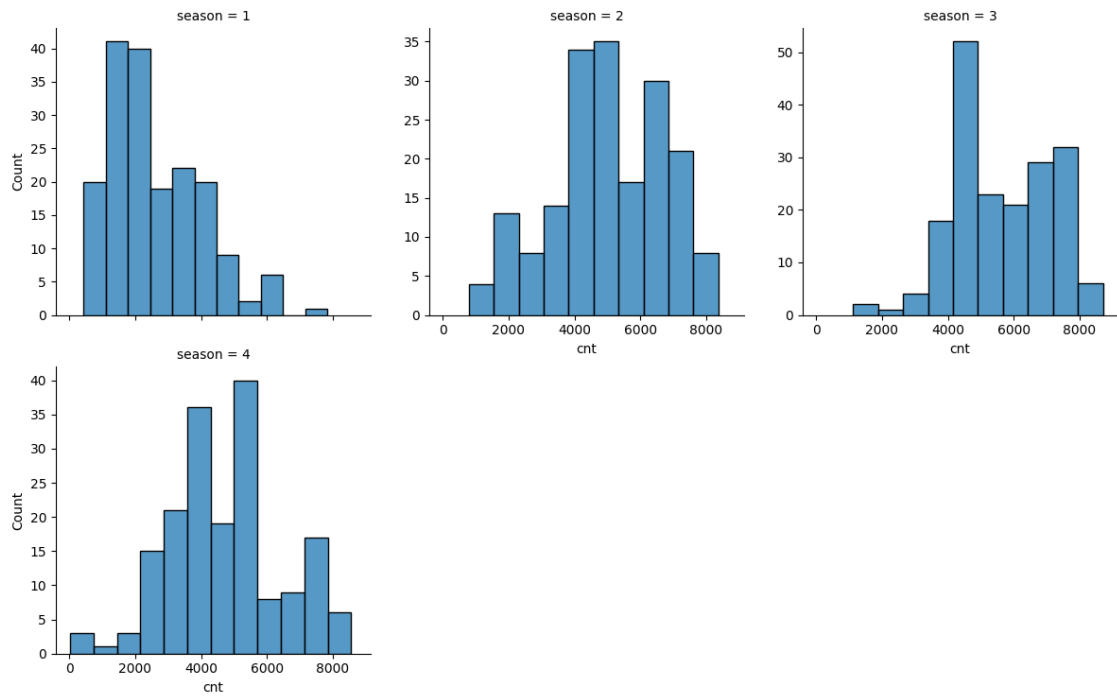
13 Viewing histograms for categorical data

```
[60]: plt.figure(figsize = (20,10))
categorical_columns = ['season', 'mnth',
      ↪'holiday', 'weekday', 'workingday', 'weathersit']
target_column = 'cnt'
for cat_column in categorical_columns:
    g = sns.FacetGrid(bike_sharing_data, col=cat_column, col_wrap=3, height=4,
      ↪sharey=False)
    g.map(sns.histplot, target_column)
    g.set_axis_labels(target_column, "Count")
    g.fig.subplots_adjust(top=0.9)
    g.fig.suptitle(f'Histograms of {target_column} for each {cat_column}')


plt.show()
```

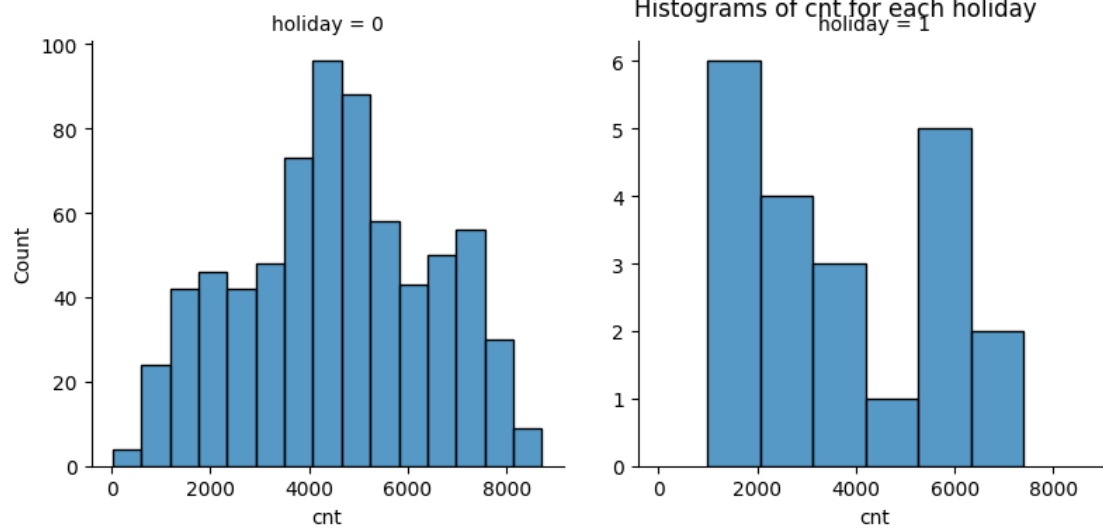
<Figure size 2000x1000 with 0 Axes>

Histograms of cnt for each season




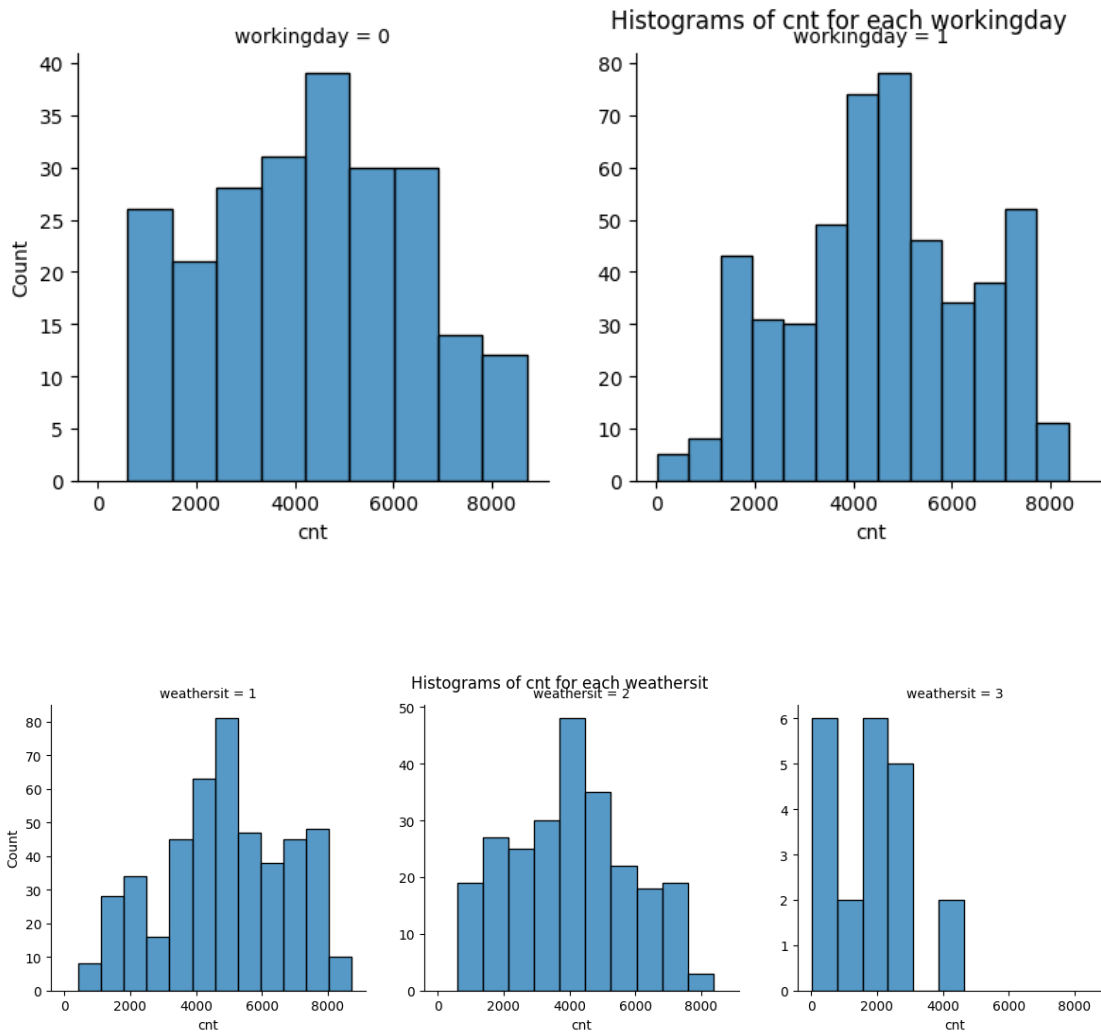
Histograms of cnt for each mnth





Histograms of cnt for each weekday





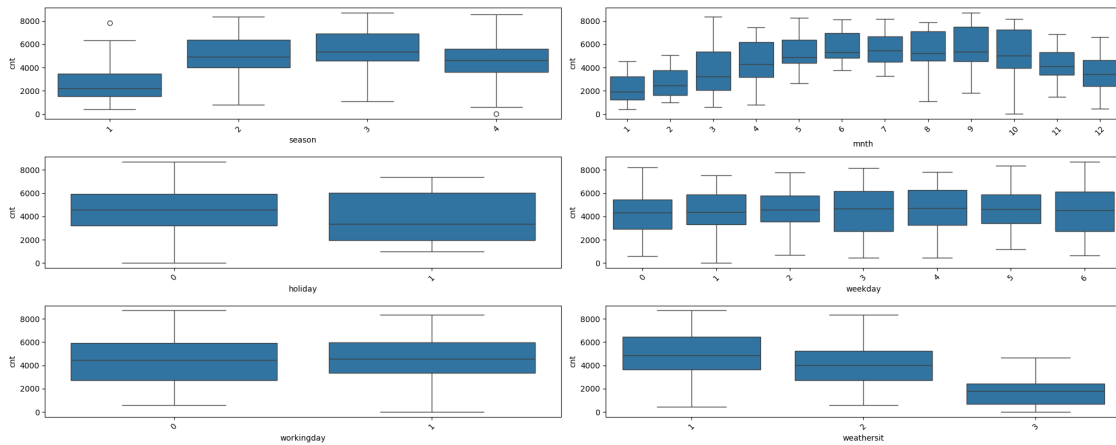
14 Understanding the distribution, using boxplots

```
[61]: plt.figure(figsize=(20, 8))
categorical_columns = ['season', 'mnth', '
    ↪ 'holiday', 'weekday', 'workingday', 'weathersit']
target_column = bike_sharing_data['cnt']

for i, col in enumerate(categorical_columns):
    plt.subplot(3, 2, i + 1)
    sns.boxplot(x=bike_sharing_data[col], y=target_column)
    plt.xlabel(col)
    plt.xticks(rotation=45)

plt.tight_layout()
```

```
plt.show()
```



15 Understanding the boxplots

season: median close to 55000, season 3 takes more booking, followed by season 2 and season 4. Out of all season 1, we have very less demand.

mnth: more booking in the months of 6, 7, 8, 9, 10 with almost the same median value. Of all considerable booking from April to November, higher demand

holiday: 3000 to 6000 bikes booked. Mean shows little above to 4000 bikes. And even on holidays, there is more booking. But median stands less than 4000.

weekday: the median lies between 4000 to 4500 and more booking on wednesdays. Very close trend each day as compared all the days in the week.

workingday: the median stands same for working day and NOT a working a day. But working day spread more even & perfect. Seems to be a good variable


weathersit: very look demand in weathersit-3. More demand in weathersit-1 then comes weather-2. Varying median in all the weathersit 1,2,3

16 viewing the correlation

16.0.1 atemp vs temp are the only two features having high correlation, indicates that they are highly linearly dependent.


16.0.2 also, all these variables are correlated with the target variable cnt.

```
[62]: plt.figure(figsize = (10,10))
sns.heatmap(bike_sharing_data.corr(), annot = True)
plt.show()
```



17 Viewing pair-plot for numeric columns

```
[63]: num_cols = bike_sharing_data[['temp', 'atemp', 'hum', 'windspeed', 'cnt']]
sns.pairplot(num_cols, diag_kind='kde')
plt.show()
```



18 Encoding of categorical columns

```
[64]: bike_sharing_data_encoded = pd.get_dummies(bike_sharing_data, drop_first=True)
bike_sharing_data_encoded.columns
```

```
[64]: Index(['yr', 'holiday', 'workingday', 'temp', 'atemp', 'hum', 'windspeed',
          'cnt', 'season_2', 'season_3', 'season_4', 'mnth_2', 'mnth_3', 'mnth_4',
          'mnth_5', 'mnth_6', 'mnth_7', 'mnth_8', 'mnth_9', 'mnth_10', 'mnth_11',
          'mnth_12', 'weekday_1', 'weekday_2', 'weekday_3', 'weekday_4',
          'weekday_5', 'weekday_6', 'weathersit_2', 'weathersit_3'],
          dtype='object')
```

19 Final list of columns, after removing unwanted ones and categorical encoding

```
[65]: bike_sharing_data_encoded.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 730 entries, 0 to 729
Data columns (total 30 columns):
#   Column                Non-Null Count  Dtype
---  -
0   yr                    730 non-null    int64
1   holiday              730 non-null    int64
2   workingday          730 non-null    int64
3   temp                 730 non-null    float64
4   atemp                730 non-null    float64
5   hum                  730 non-null    float64
6   windspeed            730 non-null    float64
7   cnt                  730 non-null    int64
8   season_2             730 non-null    bool
9   season_3             730 non-null    bool
10  season_4             730 non-null    bool
11  mnth_2               730 non-null    bool
12  mnth_3               730 non-null    bool
13  mnth_4               730 non-null    bool
14  mnth_5               730 non-null    bool
15  mnth_6               730 non-null    bool
16  mnth_7               730 non-null    bool
17  mnth_8               730 non-null    bool
18  mnth_9               730 non-null    bool
19  mnth_10              730 non-null    bool
20  mnth_11              730 non-null    bool
21  mnth_12              730 non-null    bool
22  weekday_1            730 non-null    bool
23  weekday_2            730 non-null    bool
24  weekday_3            730 non-null    bool
25  weekday_4            730 non-null    bool
26  weekday_5            730 non-null    bool
27  weekday_6            730 non-null    bool
28  weathersit_2          730 non-null    bool
29  weathersit_3          730 non-null    bool
```



```
dtypes: bool(22), float64(4), int64(4)
memory usage: 61.4 KB
```


```
[66]: bike_sharing_data.shape
```

```
[66]: (730, 12)
```

20 Re-visiting the heatmap for correlation, after encoding

20.0.1 the correlation of different variables with the target variable seems good. not dropping any column for now

```
[67]: plt.figure(figsize = (30,30))
      sns.heatmap(bike_sharing_data_encoded.corr(), annot = True)
      plt.show()
```



21 Framing x and y datasets

```
[68]: x = bike_sharing_data_encoded.drop(['cnt'], axis=1)
      y = bike_sharing_data_encoded['cnt']

      x.head()
```

```
[68]:   yr  holiday  workingday      temp      atemp      hum  windspeed  season_2 \
0     0         0              0  14.110847  18.18125  80.5833  10.749882      False
```

1	0	0	0	14.902598	17.68695	69.6087	16.652113	False
2	0	0	1	8.050924	9.47025	43.7273	16.636703	False
3	0	0	1	8.200000	10.60610	59.0435	10.739832	False
4	0	0	1	9.305237	11.46350	43.6957	12.522300	False

	season_3	season_4	...	mnth_11	mnth_12	weekday_1	weekday_2	weekday_3	\
0	False	False	...	False	False	False	False	False	
1	False	False	...	False	False	False	False	False	
2	False	False	...	False	False	True	False	False	
3	False	False	...	False	False	False	True	False	
4	False	False	...	False	False	False	False	True	

	weekday_4	weekday_5	weekday_6	weathersit_2	weathersit_3
0	False	False	True	True	False
1	False	False	False	True	False
2	False	False	False	False	False
3	False	False	False	False	False
4	False	False	False	False	False

[5 rows x 29 columns]

```
[69]: y.head()
```

```
[69]: 0    985
      1    801
      2   1349
      3   1562
      4   1600
      Name: cnt, dtype: int64
```

22 spitting the data into train and test

```
[70]: X_train, X_test, y_train, y_test = train_test_split(x, y, test_size=0.25,
      ↪random_state=101)
      print('X_train:', X_train.shape, 'X_test:', X_test.shape, 'y_train:', y_train.
      ↪shape, 'y_test:', y_test.shape)
```

```
X_train: (547, 29) X_test: (183, 29) y_train: (547,) y_test: (183,)
```

```
[71]: X_train.head()
```

```
[71]:   yr  holiday  workingday   temp   atemp   hum  windspeed  \
      182    0         0         0  30.271653  33.36540  44.4583    7.709154
      261    0         0         1  22.515847  26.48375  69.0000   10.166714
      431    1         0         1  21.627500  26.23020  56.7500   25.201627
       62    0         0         1  10.728347  12.78395  61.0417   13.624182
      473    1         0         1  20.431653  24.65230  61.2500    4.417256
```

	season_2	season_3	season_4	...	mnth_11	mnth_12	weekday_1	\
182	False	True	False	...	False	False	False	
261	False	True	False	...	False	False	True	
431	False	False	False	...	False	False	False	
62	False	False	False	...	False	False	False	
473	True	False	False	...	False	False	False	

	weekday_2	weekday_3	weekday_4	weekday_5	weekday_6	weathersit_2	\
182	False	False	False	False	True	False	
261	False	False	False	False	False	True	
431	False	False	True	False	False	False	
62	False	False	False	True	False	True	
473	False	False	True	False	False	False	

	weathersit_3
182	False
261	False
431	False
62	False
473	False

[5 rows x 29 columns]

```
[72]: y_train.head()
```

```
[72]: 182    5119
      261    4539
      431    5382
      62     1944
      473    6565
      Name: cnt, dtype: int64
```

23 Scalling X_train

```
[73]: scaler = MinMaxScaler()
      X_train[['temp', 'atemp', 'hum', 'windspeed']] = scaler.
      ↪fit_transform(X_train[['temp', 'atemp', 'hum', 'windspeed']])
      X_train.head()
```

```
[73]:   yr  holiday  workingday   temp   atemp   hum  windspeed  \
182   0         0             0  0.846320  0.772142  0.215124  0.192132
261   0         0             1  0.610610  0.591480  0.613617  0.305631
431   1         0             1  0.583612  0.584824  0.414709  1.000000
62    0         0             1  0.252371  0.231824  0.484395  0.465310
473   1         0             1  0.547268  0.543400  0.487778  0.040100
```

	season_2	season_3	season_4	...	mnth_11	mnth_12	weekday_1	\
182	False	True	False	...	False	False	False	
261	False	True	False	...	False	False	True	
431	False	False	False	...	False	False	False	
62	False	False	False	...	False	False	False	
473	True	False	False	...	False	False	False	

	weekday_2	weekday_3	weekday_4	weekday_5	weekday_6	weathersit_2	\
182	False	False	False	False	True	False	
261	False	False	False	False	False	True	
431	False	False	True	False	False	False	
62	False	False	False	True	False	True	
473	False	False	True	False	False	False	

	weathersit_3
182	False
261	False
431	False
62	False
473	False

[5 rows x 29 columns]

24 unfortunately `y_train` happens to be a series object, so, a bit lengthly code.

```
[74]: y_train_cnt = y_train.values.reshape(-1, 1)
y_train_cnt_scaled = scaler.fit_transform(y_train_cnt)
y_train_scaled = pd.Series(y_train_cnt_scaled.flatten(), name=y_train.name,
    ↪index=y_train.index)

y_train_scaled.head()
```

```
[74]: 182    0.597328
      261    0.529357
      431    0.628150
      62     0.225243
      473    0.766788
      Name: cnt, dtype: float64
```

25 Model 1 - Linear Regression model using statmodels

```
[75]: X_train['season_2'] = X_train['season_2'].map({True: 1, False: 0})
X_train['season_3'] = X_train['season_2'].map({True: 1, False: 0})
X_train['mnth_2'] = X_train['mnth_2'].map({True: 1, False: 0})
X_train['mnth_3'] = X_train['mnth_3'].map({True: 1, False: 0})
X_train['mnth_4'] = X_train['mnth_4'].map({True: 1, False: 0})
X_train['mnth_5'] = X_train['mnth_5'].map({True: 1, False: 0})
X_train['mnth_6'] = X_train['mnth_6'].map({True: 1, False: 0})
X_train['mnth_7'] = X_train['mnth_7'].map({True: 1, False: 0})
X_train['mnth_8'] = X_train['mnth_8'].map({True: 1, False: 0})
X_train['mnth_9'] = X_train['mnth_9'].map({True: 1, False: 0})
X_train['mnth_10'] = X_train['mnth_10'].map({True: 1, False: 0})
X_train['mnth_11'] = X_train['mnth_11'].map({True: 1, False: 0})
X_train['mnth_12'] = X_train['mnth_12'].map({True: 1, False: 0})
X_train['weekday_1'] = X_train['weekday_1'].map({True: 1, False: 0})
X_train['weekday_2'] = X_train['weekday_2'].map({True: 1, False: 0})
X_train['weekday_3'] = X_train['weekday_3'].map({True: 1, False: 0})
```

```
[76]: print(X_train.select_dtypes(include = 'bool'))
```

	season_4	weekday_4	weekday_5	weekday_6	weathersit_2	weathersit_3
182	False	False	False	True	False	False
261	False	False	False	False	True	False
431	False	True	False	False	False	False
62	False	False	True	False	True	False
473	False	True	False	False	False	False
..
75	False	True	False	False	False	False
599	False	True	False	False	False	False
575	False	False	False	False	False	False
337	True	False	False	False	False	False
523	False	False	True	False	False	False

[547 rows x 6 columns]

```
[77]: print(y_train_scaled.dtypes)
```

float64

```
[78]: X_train_sm = sm.add_constant(X_train)
lr_1 = sm.OLS(np.asarray(y_train_scaled), X_train_sm)
lr_1_model = lr_1.fit()

lr_1_model.params
```

ValueError

Traceback (most recent call last)

Cell In[78], line 2

```
1 X_train_sm = sm.add_constant(X_train)
----> 2 lr_1 = sm.OLS(np.asarray(y_train_scaled), X_train_sm)
3 lr_1_model = lr_1.fit()
5 lr_1_model.params
```

File c:

```
↪ \Users\sndur\AppData\Local\Programs\Python\Python311\Lib\site-packages\statsmodels\regress
↪ py:923, in OLS.__init__(self, endog, exog, missing, hasconst, **kwargs)
920     msg = ("Weights are not supported in OLS and will be ignored"
921           "An exception will be raised in the next version.")
922     warnings.warn(msg, ValueWarning)
--> 923 super(OLS, self).__init__(endog, exog, missing=missing,
924                             hasconst=hasconst, **kwargs)
925 if "weights" in self._init_keys:
926     self._init_keys.remove("weights")
```

File c:

```
↪ \Users\sndur\AppData\Local\Programs\Python\Python311\Lib\site-packages\statsmodels\regress
↪ py:748, in WLS.__init__(self, endog, exog, weights, missing, hasconst,
↪ **kwargs)
746 else:
747     weights = weights.squeeze()
--> 748 super(WLS, self).__init__(endog, exog, missing=missing,
749                             weights=weights, hasconst=hasconst, **kwargs)
750 nobs = self.exog.shape[0]
751 weights = self.weights
```

File c:

```
↪ \Users\sndur\AppData\Local\Programs\Python\Python311\Lib\site-packages\statsmodels\regress
↪ py:202, in RegressionModel.__init__(self, endog, exog, **kwargs)
201 def __init__(self, endog, exog, **kwargs):
--> 202     super(RegressionModel, self).__init__(endog, exog, **kwargs)
203     self.pinv_wexog: Float64Array | None = None
204     self._data_attr.extend(['pinv_wexog', 'wendog', 'wexog', 'weights'])
```

File c:

```
↪ \Users\sndur\AppData\Local\Programs\Python\Python311\Lib\site-packages\statsmodels\base\mo
↪ py:270, in LikelihoodModel.__init__(self, endog, exog, **kwargs)
269 def __init__(self, endog, exog=None, **kwargs):
--> 270     super().__init__(endog, exog, **kwargs)
271     self.initialize()
```

File c:

```
↪ \Users\sndur\AppData\Local\Programs\Python\Python311\Lib\site-packages\statsmodels\base\mo
↪ py:95, in Model.__init__(self, endog, exog, **kwargs)
93 missing = kwargs.pop('missing', 'none')
94 hasconst = kwargs.pop('hasconst', None)
--> 95 self.data = self._handle_data(endog, exog, missing, hasconst,
```

```

96         **kwargs)
97 self.k_constant = self.data.k_constant
98 self.exog = self.data.exog

```

File c:

```

↪ \Users\sndur\AppData\Local\Programs\Python\Python311\Lib\site-packages\statsmodels\base\model
↪ py:135, in Model._handle_data(self, endog, exog, missing, hasconst, **kwargs)
    134 def _handle_data(self, endog, exog, missing, hasconst, **kwargs):
--> 135     data = handle_data(endog, exog, missing, hasconst, **kwargs)
    136     # kwargs arrays could have changed, easier to just attach here
    137     for key in kwargs:

```

File c:

```

↪ \Users\sndur\AppData\Local\Programs\Python\Python311\Lib\site-packages\statsmodels\base\data
↪ py:675, in handle_data(endog, exog, missing, hasconst, **kwargs)
    672     exog = np.asarray(exog)
    674 klass = handle_data_class_factory(endog, exog)
--> 675 return klass(endog, exog=exog, missing=missing, hasconst=hasconst,
    676                 **kwargs)

```

File c:

```

↪ \Users\sndur\AppData\Local\Programs\Python\Python311\Lib\site-packages\statsmodels\base\data
↪ py:84, in ModelData.__init__(self, endog, exog, missing, hasconst, **kwargs)
    82     self.orig_endog = endog
    83     self.orig_exog = exog
---> 84     self.endog, self.exog = self._convert_endog_exog(endog, exog)
    86 self.const_idx = None
    87 self.k_constant = 0

```

File c:

```

↪ \Users\sndur\AppData\Local\Programs\Python\Python311\Lib\site-packages\statsmodels\base\data
↪ py:509, in PandasData._convert_endog_exog(self, endog, exog)
    507 exog = exog if exog is None else np.asarray(exog)
    508 if endog.dtype == object or exog is not None and exog.dtype == object:
--> 509     raise ValueError("Pandas data cast to numpy dtype of object. "
    510                       "Check input data with np.asarray(data).")
    511 return super(PandasData, self)._convert_endog_exog(endog, exog)

```

ValueError: Pandas data cast to numpy dtype of object. Check input data with np
↪ asarray(data).

26 In the below summary table:

26.0.1 R-squared: 0.846, means that approximately 84.6% of the variability in the dependent variable is explained by the independent variables in your model. This is a relatively high R-squared value, indicating that the model is quite effective in explaining and predicting the variation in the dependent variable

26.0.2 F-statistic (101.3) suggests that the model, as a whole, is more likely to be statistically significant

26.0.3 Prob (F-statistic): 4.33e-190 suggests that the overall model is statistically significant. More specifically

```
[ ]: lr_1_model.summary()
```

```
[ ]: <class 'statsmodels.iolib.summary.Summary'>
      """
```

```

                        OLS Regression Results
=====
Dep. Variable:          cnt      R-squared:                0.846
Model:                  OLS      Adj. R-squared:           0.837
Method:                 Least Squares      F-statistic:        101.3
Date:                  Wed, 13 Dec 2023      Prob (F-statistic):    4.33e-190
Time:                  11:42:30      Log-Likelihood:       546.48
No. Observations:      547      AIC:                  -1035.
Df Residuals:          518      BIC:                  -910.1
Df Model:               28
Covariance Type:       nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	0.1416	0.023	6.072	0.000	0.096	0.187
yr	0.2368	0.008	29.140	0.000	0.221	0.253
holiday	-0.0579	0.024	-2.442	0.015	-0.104	-0.011
workingday	0.0373	0.009	3.997	0.000	0.019	0.056
temp	0.0392	0.258	0.152	0.879	-0.468	0.546
atemp	0.4125	0.265	1.556	0.120	-0.108	0.933
hum	-0.1196	0.027	-4.508	0.000	-0.172	-0.067
windspeed	-0.1143	0.020	-5.808	0.000	-0.153	-0.076
season_2	0.0902	0.025	3.657	0.000	0.042	0.139
season_3	0.0740	0.030	2.505	0.013	0.016	0.132
season_4	0.1863	0.026	7.110	0.000	0.135	0.238
mnth_2	0.0180	0.020	0.917	0.360	-0.021	0.056
mnth_3	0.0731	0.022	3.317	0.001	0.030	0.116
mnth_4	0.0596	0.034	1.760	0.079	-0.007	0.126
mnth_5	0.0967	0.037	2.645	0.008	0.025	0.169
mnth_6	0.0651	0.039	1.662	0.097	-0.012	0.142
mnth_7	0.0168	0.043	0.386	0.700	-0.069	0.102
mnth_8	0.0683	0.042	1.626	0.105	-0.014	0.151

mnth_9	0.1277	0.037	3.434	0.001	0.055	0.201
mnth_10	0.0489	0.034	1.421	0.156	-0.019	0.117
mnth_11	-0.0254	0.033	-0.770	0.441	-0.090	0.039
mnth_12	-0.0174	0.027	-0.647	0.518	-0.070	0.035
weekday_1	-0.0194	0.010	-1.979	0.048	-0.039	-0.000
weekday_2	-0.0086	0.011	-0.776	0.438	-0.031	0.013
weekday_3	-0.0020	0.011	-0.176	0.860	-0.024	0.020
weekday_4	0.0049	0.010	0.471	0.638	-0.015	0.025
weekday_5	0.0046	0.011	0.424	0.672	-0.017	0.026
weekday_6	0.0503	0.015	3.459	0.001	0.022	0.079
weathersit_2	-0.0529	0.011	-4.789	0.000	-0.075	-0.031
weathersit_3	-0.2325	0.029	-7.977	0.000	-0.290	-0.175

Omnibus:	97.916	Durbin-Watson:	1.886
Prob(Omnibus):	0.000	Jarque-Bera (JB):	251.917
Skew:	-0.901	Prob(JB):	1.98e-55
Kurtosis:	5.794	Cond. No.	2.25e+15

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The smallest eigenvalue is 3.59e-28. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

"""

27 Model 1 -> find VIF (i) = 1 / (1 - rsquare)

```
[ ]: vif = pd.DataFrame()
vif['Features'] = X_train.columns
vif['VIF'] = [variance_inflation_factor(X_train.values,i) for i in
↳range(X_train.shape[1])]
vif['VIF'] = round(vif['VIF'],2)
vif = vif.sort_values(by = 'VIF', ascending=False)
vif
```

```
[ ]:
24    weekday_4    inf
2    workingday    inf
23    weekday_3    inf
22    weekday_2    inf
21    weekday_1    inf
25    weekday_5    inf
1     holiday     inf
3         temp 1506.52
4         atemp 1434.59
```

8	season_3	14.67
5	hum	11.81
15	mnth_7	10.59
9	season_4	10.57
7	season_2	10.02
16	mnth_8	9.26
14	mnth_6	8.79
13	mnth_5	7.34
17	mnth_9	7.25
18	mnth_10	6.20
12	mnth_4	6.12
19	mnth_11	5.31
6	windspeed	4.40
20	mnth_12	3.69
11	mnth_3	2.78
27	weathersit_2	2.60
0	yr	2.13
26	weekday_6	1.89
10	mnth_2	1.81
28	weathersit_3	1.46

28 Proceeding with the below interpretation

28.1 Features with High p-value and High VIF:

28.1.1 Interpretation: These features are both statistically insignificant and contribute to multicollinearity.

28.1.2 Action: Consider removing these features from the model. They don't provide significant explanatory power, and their presence may lead to multicollinearity issues.

28.2 Features with High p-value and Low VIF:

28.2.1 Interpretation: These features are statistically insignificant but don't contribute significantly to multicollinearity.

28.2.2 Action: Evaluate the importance of these features in the context of your problem. If they are not essential, consider removing them. If they are important, you may want to explore other transformations or interactions.

28.3 Features with Low p-value and High VIF:

28.3.1 Interpretation: These features are statistically significant, but their presence may contribute to multicollinearity.

28.3.2 Action: Assess the importance of these features. If they are crucial for your analysis, you might consider keeping them but explore options to address multicollinearity, such as feature engineering, regularization, or removing one of the correlated variables.

28.4 Features with Low p-value and Low VIF:

28.4.1 Interpretation: These features are statistically significant and do not exhibit high multicollinearity.

28.4.2 Action: These are generally desirable features to include in your model. They contribute meaningfully to the prediction without introducing multicollinearity concerns.

29 Model 1 - Removing the features

`'weekday_4','weekday_3','weekday_2','weekday_1','weekday_5','weekday_6','holiday','temp'`

really sure why these weekday_n columns are showing infinite 'inf' as VIF. They should have some value. Followed all basics in the coding and these are just categorical encoded columns. Anyway deleting them.

```
[ ]: X_train_sm.  
      ↪ drop(['weekday_4', 'weekday_3', 'weekday_2', 'weekday_1', 'weekday_5', 'weekday_6', 'holiday', 'at  
      ↪ axis=1, inplace=True)  
X_train_sm.columns
```

```
[ ]: Index(['const', 'yr', 'workingday', 'temp', 'hum', 'windspeed', 'season_2',  
          'season_3', 'season_4', 'mnth_2', 'mnth_3', 'mnth_4', 'mnth_5',
```

```

'mnth_6', 'mnt_7', 'mnt_8', 'mnt_9', 'mnt_10', 'mnt_11', 'mnt_12',
'weathersit_2', 'weathersit_3'],
dtype='object')

```

30 Model 2

```

[ ]: #X_train_sm = sm.add_constant(X_train)
lr_2 = sm.OLS(y_train_scaled, X_train_sm)
lr_2_model = lr_2.fit()

lr_2_model.params

```

```

[ ]: const          0.161548
yr              0.236744
workingday      0.015476
temp            0.427190
hum            -0.119108
windspeed      -0.124409
season_2        0.089848
season_3        0.073186
season_4        0.189928
mnt_2           0.019817
mnt_3           0.080483
mnt_4           0.067760
mnt_5           0.100508
mnt_6           0.070580
mnt_7           0.022345
mnt_8           0.072521
mnt_9           0.129320
mnt_10          0.055612
mnt_11          -0.024016
mnt_12          -0.010312
weathersit_2     -0.051371
weathersit_3     -0.229131
dtype: float64

```

```

[ ]: lr_2_model.summary()

```

```

[ ]: <class 'statsmodels.iolib.summary.Summary'>
"""

```

```

                                OLS Regression Results
=====
Dep. Variable:                  cnt      R-squared:                0.836
Model:                            OLS      Adj. R-squared:            0.829
Method:                 Least Squares      F-statistic:                127.1
Date:                 Wed, 13 Dec 2023      Prob (F-statistic):        2.63e-190
Time:                  11:42:30      Log-Likelihood:            529.36

```

No. Observations: 547 AIC: -1015.
Df Residuals: 525 BIC: -920.0
Df Model: 21
Covariance Type: nonrobust

	coef	std err	t	P> t	[0.025	0.975]
const	0.1615	0.023	7.177	0.000	0.117	0.206
yr	0.2367	0.008	28.497	0.000	0.220	0.253
workingday	0.0155	0.009	1.780	0.076	-0.002	0.033
temp	0.4272	0.047	9.149	0.000	0.335	0.519
hum	-0.1191	0.027	-4.454	0.000	-0.172	-0.067
windspeed	-0.1244	0.019	-6.508	0.000	-0.162	-0.087
season_2	0.0898	0.025	3.556	0.000	0.040	0.139
season_3	0.0732	0.030	2.421	0.016	0.014	0.133
season_4	0.1899	0.027	7.108	0.000	0.137	0.242
mnth_2	0.0198	0.020	0.989	0.323	-0.020	0.059
mnth_3	0.0805	0.022	3.587	0.000	0.036	0.125
mnth_4	0.0678	0.035	1.958	0.051	-0.000	0.136
mnth_5	0.1005	0.037	2.691	0.007	0.027	0.174
mnth_6	0.0706	0.040	1.781	0.076	-0.007	0.148
mnth_7	0.0223	0.044	0.507	0.612	-0.064	0.109
mnth_8	0.0725	0.043	1.703	0.089	-0.011	0.156
mnth_9	0.1293	0.038	3.427	0.001	0.055	0.203
mnth_10	0.0556	0.035	1.583	0.114	-0.013	0.125
mnth_11	-0.0240	0.034	-0.713	0.476	-0.090	0.042
mnth_12	-0.0103	0.027	-0.377	0.707	-0.064	0.043
weathersit_2	-0.0514	0.011	-4.603	0.000	-0.073	-0.029
weathersit_3	-0.2291	0.029	-7.859	0.000	-0.286	-0.172
Omnibus:	85.499	Durbin-Watson:	1.904			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	226.252			
Skew:	-0.782	Prob(JB):	7.41e-50			
Kurtosis:	5.735	Cond. No.	43.6			

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

"""

31 VIF after Model 2

```
[ ]: vif = pd.DataFrame()
vif['Features'] = X_train_sm.columns
```

```
vif['VIF'] = [variance_inflation_factor(X_train_sm.values,i) for i in   

↳range(X_train_sm.shape[1])]
vif['VIF'] = round(vif['VIF'],2)
vif = vif.sort_values(by = 'VIF', ascending=False)
vif
```

```
[ ]:      Features      VIF
0      const  31.47
7    season_3  10.86
14    mnth_7   9.48
15    mnth_8   8.33
8    season_4   7.99
13    mnth_6   7.81
6    season_2   7.48
3      temp   7.13
12    mnth_5   6.67
16    mnth_9   6.54
17    mnth_10   5.67
11    mnth_4   5.62
18    mnth_11   5.00
19    mnth_12   3.51
10    mnth_3   2.60
4      hum    2.29
9    mnth_2    1.85
20 weathersit_2  1.71
21 weathersit_3  1.41
5    windspeed  1.26
1      yr      1.07
2    workingday  1.02
```

32 Model 2 -Removing features based on the explained criteria of p-value and vif

```
[ ]: X_train_sm.drop(['mnth_6', 'mnth_4', 'mnth_6', 'mnth_3', 'mnth_8', 'mnth_9',   

↳'mnth_10', 'mnth_12'], axis=1, inplace=True)
X_train_sm.columns
```

```
[ ]: Index(['const', 'yr', 'workingday', 'temp', 'hum', 'windspeed', 'season_2',   

'season_3', 'season_4', 'mnth_2', 'mnth_5', 'mnth_7', 'mnth_11',   

'weathersit_2', 'weathersit_3'],   

dtype='object')
```

33 Model 3

```
[ ]: #X_train_sm = sm.add_constant(X_train)
lr_3 = sm.OLS(y_train_scaled, X_train_sm)
lr_3_model = lr_3.fit()

lr_3_model.params
```

```
[ ]: const          0.159836
     yr             0.234738
     workingday     0.016842
     temp           0.501077
     hum            -0.113947
     windspeed      -0.115592
     season_2       0.112901
     season_3       0.102954
     season_4       0.190347
     mnth_2         -0.007744
     mnth_5         0.022623
     mnth_7         -0.075229
     mnth_11        -0.057754
     weathersit_2    -0.049597
     weathersit_3    -0.229116
     dtype: float64
```

```
[ ]: lr_3_model.summary()
```

```
[ ]: <class 'statsmodels.iolib.summary.Summary'>
     """
```

```

                        OLS Regression Results
=====
Dep. Variable:          cnt      R-squared:                0.823
Model:                  OLS      Adj. R-squared:           0.818
Method:                 Least Squares      F-statistic:          176.7
Date:                   Wed, 13 Dec 2023    Prob (F-statistic):      1.49e-189
Time:                   11:42:30           Log-Likelihood:         509.12
No. Observations:       547              AIC:                  -988.2
Df Residuals:           532              BIC:                  -923.7
Df Model:               14
Covariance Type:        nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	0.1598	0.022	7.407	0.000	0.117	0.202
yr	0.2347	0.008	27.784	0.000	0.218	0.251
workingday	0.0168	0.009	1.884	0.060	-0.001	0.034
temp	0.5011	0.035	14.415	0.000	0.433	0.569

hum	-0.1139	0.027	-4.286	0.000	-0.166	-0.062
windspeed	-0.1156	0.019	-5.958	0.000	-0.154	-0.077
season_2	0.1129	0.017	6.614	0.000	0.079	0.146
season_3	0.1030	0.022	4.782	0.000	0.061	0.145
season_4	0.1903	0.016	12.073	0.000	0.159	0.221
mnth_2	-0.0077	0.018	-0.437	0.662	-0.043	0.027
mnth_5	0.0226	0.018	1.257	0.209	-0.013	0.058
mnth_7	-0.0752	0.018	-4.246	0.000	-0.110	-0.040
mnth_11	-0.0578	0.019	-3.119	0.002	-0.094	-0.021
weathersit_2	-0.0496	0.011	-4.343	0.000	-0.072	-0.027
weathersit_3	-0.2291	0.030	-7.663	0.000	-0.288	-0.170

```
=====
Omnibus:                65.758    Durbin-Watson:                1.894
Prob(Omnibus):           0.000    Jarque-Bera (JB):           160.137
Skew:                    -0.632    Prob(JB):                   1.69e-35
Kurtosis:                5.330    Cond. No.                   17.6
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

"""

34 VIF after Model 3

```
[ ]: vif = pd.DataFrame()
vif['Features'] = X_train_sm.columns
vif['VIF'] = [variance_inflation_factor(X_train_sm.values,i) for i in
↳range(X_train_sm.shape[1])]
vif['VIF'] = round(vif['VIF'],2)
vif = vif.sort_values(by = 'VIF', ascending=False)
vif
```

```
[ ]:      Features    VIF
0      const    27.22
7    season_3     5.18
3      temp     3.72
6    season_2     3.21
8    season_4     2.62
4      hum     2.13
13 weathersit_2     1.69
10     mnth_5     1.46
11     mnth_7     1.44
12     mnth_11     1.42
14 weathersit_3     1.39
9      mnth_2     1.36
```

```

5      windspeed    1.23
1              yr    1.04
2      workingday    1.02

```

----- RFE -----

35 Model 3 - Removing features based on the explained criteria of p-value and vif

35.0.1 'hum', 'season_3'

```

[ ]: X_train_sm.drop(['hum', 'season_3'], axis=1, inplace=True)
      X_train_sm.columns

[ ]: Index(['const', 'yr', 'workingday', 'temp', 'windspeed', 'season_2',
           'season_4', 'mnth_2', 'mnth_5', 'mnth_7', 'mnth_11', 'weathersit_2',
           'weathersit_3'],
          dtype='object')

```

36 Model 4

```

[ ]: #X_train_sm = sm.add_constant(X_train)
      lr_4 = sm.OLS(y_train_scaled, X_train_sm)
      lr_4_model = lr_4.fit()

      lr_4_model.params

```

```

[ ]: const          0.099902
      yr            0.236613
      workingday    0.017476
      temp          0.600506
      windspeed     -0.094492
      season_2      0.063354
      season_4      0.138822
      mnth_2        -0.025982
      mnth_5         0.000188
      mnth_7        -0.054405
      mnth_11       -0.044890
      weathersit_2   -0.076087
      weathersit_3   -0.284441
      dtype: float64

```

```

[ ]: lr_4_model.summary()

```

```

[ ]: <class 'statsmodels.iolib.summary.Summary'>
      """
                                     OLS Regression Results

```

```

=====
Dep. Variable:          cnt    R-squared:                0.809
Model:                  OLS    Adj. R-squared:           0.804
Method:                 Least Squares    F-statistic:          188.0
Date:                  Wed, 13 Dec 2023    Prob (F-statistic):    7.64e-183
Time:                  11:42:30    Log-Likelihood:        487.74
No. Observations:      547    AIC:                  -949.5
Df Residuals:          534    BIC:                  -893.5
Df Model:              12
Covariance Type:       nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	0.0999	0.019	5.239	0.000	0.062	0.137
yr	0.2366	0.009	27.404	0.000	0.220	0.254
workingday	0.0175	0.009	1.886	0.060	-0.001	0.036
temp	0.6005	0.023	25.609	0.000	0.554	0.647
windspeed	-0.0945	0.019	-4.909	0.000	-0.132	-0.057
season_2	0.0634	0.013	4.905	0.000	0.038	0.089
season_4	0.1388	0.013	10.499	0.000	0.113	0.165
mnth_2	-0.0260	0.018	-1.465	0.143	-0.061	0.009
mnth_5	0.0002	0.018	0.010	0.992	-0.036	0.036
mnth_7	-0.0544	0.018	-3.007	0.003	-0.090	-0.019
mnth_11	-0.0449	0.019	-2.355	0.019	-0.082	-0.007
weathersit_2	-0.0761	0.009	-8.077	0.000	-0.095	-0.058
weathersit_3	-0.2844	0.027	-10.557	0.000	-0.337	-0.232

```

=====
Omnibus:                53.239    Durbin-Watson:           1.913
Prob(Omnibus):          0.000    Jarque-Bera (JB):        107.001
Skew:                   -0.576    Prob(JB):                5.82e-24
Kurtosis:               4.835    Cond. No.                 11.4
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

"""

37 VIF after Model 4

```

[ ]: vif = pd.DataFrame()
vif['Features'] = X_train_sm.columns
vif['VIF'] = [variance_inflation_factor(X_train_sm.values,i) for i in
↳range(X_train_sm.shape[1])]
vif['VIF'] = round(vif['VIF'],2)
vif = vif.sort_values(by = 'VIF', ascending=False)

```

```
vif
```

```
[ ]:      Features      VIF
0         const    19.73
5      season_2     1.71
6      season_4     1.71
3         temp     1.57
8       mnth_5     1.40
10      mnth_11     1.40
9       mnth_7     1.39
7       mnth_2     1.26
4     windspeed     1.12
11  weathersit_2     1.07
12  weathersit_3     1.05
1             yr     1.01
2    workingday     1.01
```

38 Parameters after Model 4, we can remove `mnth_2` and `mnth_5`, as the p-value is more and VIF is low, these two become insignificant

```
[ ]: X_train_sm.drop(['mnth_2', 'mnth_5'], axis=1, inplace=True)
X_train_sm.columns
```

```
[ ]: Index(['const', 'yr', 'workingday', 'temp', 'windspeed', 'season_2',
        'season_4', 'mnth_7', 'mnth_11', 'weathersit_2', 'weathersit_3'],
        dtype='object')
```

39 Model 5

```
[ ]: #X_train_sm = sm.add_constant(X_train)
lr_5 = sm.OLS(y_train_scaled, X_train_sm)
lr_5_model = lr_5.fit()

lr_5_model.params
```

```
[ ]: const          0.091109
yr              0.236081
workingday      0.016914
temp            0.610654
windspeed      -0.096030
season_2        0.067240
season_4        0.143748
mnth_7         -0.053292
mnth_11        -0.043904
```

```
weathersit_2    -0.075563
weathersit_3    -0.283563
dtype: float64
```

```
[ ]: lr_5_model.summary()
```

```
[ ]: <class 'statsmodels.iolib.summary.Summary'>
      """
```

```

                                OLS Regression Results
=====
Dep. Variable:                  cnt    R-squared:                  0.808
Model:                        OLS    Adj. R-squared:              0.804
Method:                    Least Squares    F-statistic:                225.4
Date:                Wed, 13 Dec 2023    Prob (F-statistic):        9.85e-185
Time:                        11:42:31    Log-Likelihood:            486.64
No. Observations:                547    AIC:                       -951.3
Df Residuals:                    536    BIC:                       -903.9
Df Model:                        10
Covariance Type:                nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	0.0911	0.018	5.038	0.000	0.056	0.127
yr	0.2361	0.009	27.363	0.000	0.219	0.253
workingday	0.0169	0.009	1.827	0.068	-0.001	0.035
temp	0.6107	0.022	27.408	0.000	0.567	0.654
windspeed	-0.0960	0.019	-5.009	0.000	-0.134	-0.058
season_2	0.0672	0.011	6.034	0.000	0.045	0.089
season_4	0.1437	0.013	11.238	0.000	0.119	0.169
mnth_7	-0.0533	0.018	-2.951	0.003	-0.089	-0.018
mnth_11	-0.0439	0.019	-2.304	0.022	-0.081	-0.006
weathersit_2	-0.0756	0.009	-8.033	0.000	-0.094	-0.057
weathersit_3	-0.2836	0.027	-10.527	0.000	-0.336	-0.231

```

=====
Omnibus:                    50.158    Durbin-Watson:              1.905
Prob(Omnibus):                0.000    Jarque-Bera (JB):           100.181
Skew:                        -0.548    Prob(JB):                   1.76e-22
Kurtosis:                    4.787    Cond. No.                   10.8
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
"""
```

40 VIF after Model 5

```
[ ]: vif = pd.DataFrame()
vif['Features'] = X_train_sm.columns
vif['VIF'] = [variance_inflation_factor(X_train_sm.values,i) for i in
↳range(X_train_sm.shape[1])]
vif['VIF'] = round(vif['VIF'],2)
vif = vif.sort_values(by = 'VIF', ascending=False)
vif
```

```
[ ]:      Features      VIF
0      const    17.74
6    season_4     1.60
3      temp     1.42
8    mnth_11     1.40
7    mnth_7     1.39
5    season_2     1.27
4    windspeed     1.11
9  weathersit_2     1.06
10 weathersit_3     1.05
1      yr         1.01
2    workingday     1.01
```


41 Residual Analysis

```
[ ]: y_train_pred = lr_5_model.predict(X_train_sm)
y_train_pred
```

```
[ ]: 182    0.536175
261    0.375980
431    0.604458
62     0.141887
473    0.741684
...
75     0.332270
599    0.819176
575    0.772007
337    0.432388
523    0.823930
Length: 547, dtype: float64
```

42 Visualizing the residuals - not a 100% normally distributed one :-)

```
[ ]: res = y_train - y_train_pred
x = sns.histplot(res)
plt.show()
```



43 Predicting the Model on test data

```
[ ]: scaler = MinMaxScaler()
X_test[['temp', 'atemp', 'hum', 'windspeed']] = scaler.
    ↪fit_transform(X_test[['temp', 'atemp', 'hum', 'windspeed']])
X_test.head()
```

```
[ ]:      yr  holiday  workingday      temp      atemp      hum  windspeed  \
299    0         0           1  0.502885  0.478679  0.813203  0.447619
723    1         1           0  0.261711  0.249188  0.686333  0.357702
647    1         0           1  0.562495  0.543746  0.517546  0.416832
520    1         0           1  0.598484  0.575859  0.489130  0.484197
114    0         0           1  0.687336  0.644471  0.754342  0.409128
```

	season_2	season_3	season_4	...	mnth_11	mnth_12	weekday_1	\
299	0	0	1	...	0	0	0	
723	0	0	0	...	0	1	0	
647	0	0	1	...	0	0	0	
520	1	0	0	...	0	0	0	
114	1	0	0	...	0	0	1	

	weekday_2	weekday_3	weekday_4	weekday_5	weekday_6	weathersit_2	\
299	0	0	1	0	0	1	
723	1	0	0	0	0	1	
647	0	1	0	0	0	0	
520	1	0	0	0	0	1	
114	0	0	0	0	0	0	

	weathersit_3
299	0
723	0
647	0
520	0
114	0

[5 rows x 29 columns]

```
[ ]: y_test_cnt = y_test.values.reshape(-1, 1)
      y_test_cnt_scaled = scaler.fit_transform(y_test_cnt)
      y_test_scaled = pd.Series(y_test_cnt_scaled.flatten(), name=y_test.name,
      ↪ index=y_test.index)
```

```
[ ]: X_test.drop(['holiday', 'hum', 'season_3', 'atemp',
      ↪ 'weekday_4', 'weekday_3', 'weekday_2', 'weekday_1', 'weekday_5', 'weekday_6',
      ↪ 'mnth_6', 'mnth_4', 'mnth_3', 'mnth_2', 'mnth_5', 'mnth_8',
      ↪ 'mnth_9', 'mnth_10', 'mnth_12'], axis=1, inplace=True)
      X_test.columns
```

```
[ ]: Index(['yr', 'workingday', 'temp', 'windspeed', 'season_2', 'season_4',
      ↪ 'mnth_7', 'mnth_11', 'weathersit_2', 'weathersit_3'],
      ↪ dtype='object')
```

```
[ ]: X_test_sm = sm.add_constant(X_test)
      y_test_pred = lr_5_model.predict(X_test_sm)
      y_test_pred
```

```
[ ]: 299    0.440311
      723    0.377091
      647    0.791312
      520    0.654749
      114    0.555698
```



```

...
227    0.550901
424    0.612566
80     0.407693
98     0.261043
467    0.665359
Length: 183, dtype: float64

```

44 Evaluating the Model after test run, using `r2_score`

45 `r2score` on test = 0.808 is a good score

46 And variable 'temp' has the highest co-efficient of 0.6017

```
[ ]: print(r2_score(y_true=y_test_scaled, y_pred=y_test_pred))

0.808456818436182
```

47 Using RFE of sklearn

```
[ ]: from sklearn.feature_selection import RFE
    from sklearn.linear_model import LinearRegression

[ ]: print(X_train.shape, y_train.shape)

(547, 29) (547,)

[ ]: lm = LinearRegression()
    lm.fit(X_train, y_train)
    rfe = RFE(lm, step=15)           # top ten features
    rfe = rfe.fit(X_train, y_train)
```

```
[ ]: list(zip(X_train.columns, rfe.support_, rfe.ranking_))
```

```
[ ]: [('yr', True, 1),
      ('holiday', True, 1),
      ('workingday', True, 1),
      ('temp', False, 2),
      ('atemp', True, 1),
      ('hum', True, 1),
      ('windspeed', True, 1),
      ('season_2', False, 2),
      ('season_3', False, 2),
      ('season_4', True, 1),
      ('mnth_2', False, 2),
      ('mnth_3', False, 2),
```

```
( 'mnth_4', False, 2),
( 'mnth_5', False, 2),
( 'mnth_6', False, 2),
( 'mnth_7', False, 2),
( 'mnth_8', False, 2),
( 'mnth_9', True, 1),
( 'mnth_10', False, 2),
( 'mnth_11', False, 2),
( 'mnth_12', False, 2),
( 'weekday_1', True, 1),
( 'weekday_2', True, 1),
( 'weekday_3', True, 1),
( 'weekday_4', True, 1),
( 'weekday_5', True, 1),
( 'weekday_6', False, 2),
( 'weathersit_2', False, 2),
( 'weathersit_3', True, 1)]
```

48 Support cols

```
[ ]: X_train_support_cols = X_train.columns[rfe.support_]
X_train_support_cols
```

```
[ ]: Index(['yr', 'holiday', 'workingday', 'atemp', 'hum', 'windspeed', 'season_4',
          'mnth_9', 'weekday_1', 'weekday_2', 'weekday_3', 'weekday_4',
          'weekday_5', 'weathersit_3'],
          dtype='object')
```

49 Not supported cols

```
[ ]: X_train_support_cols = X_train.columns[~rfe.support_]
X_train_support_cols
```

```
[ ]: Index(['temp', 'season_2', 'season_3', 'mnth_2', 'mnth_3', 'mnth_4', 'mnth_5',
          'mnth_6', 'mnth_7', 'mnth_8', 'mnth_10', 'mnth_11', 'mnth_12',
          'weekday_6', 'weathersit_2'],
          dtype='object')
```

50 We see the best supported features (from sklearn) are well included in the final features Model_5 (using statsmodel).

51 Further if we run the Model and performs the same analysis using VIF and p-value, it is expected to the last 11 features